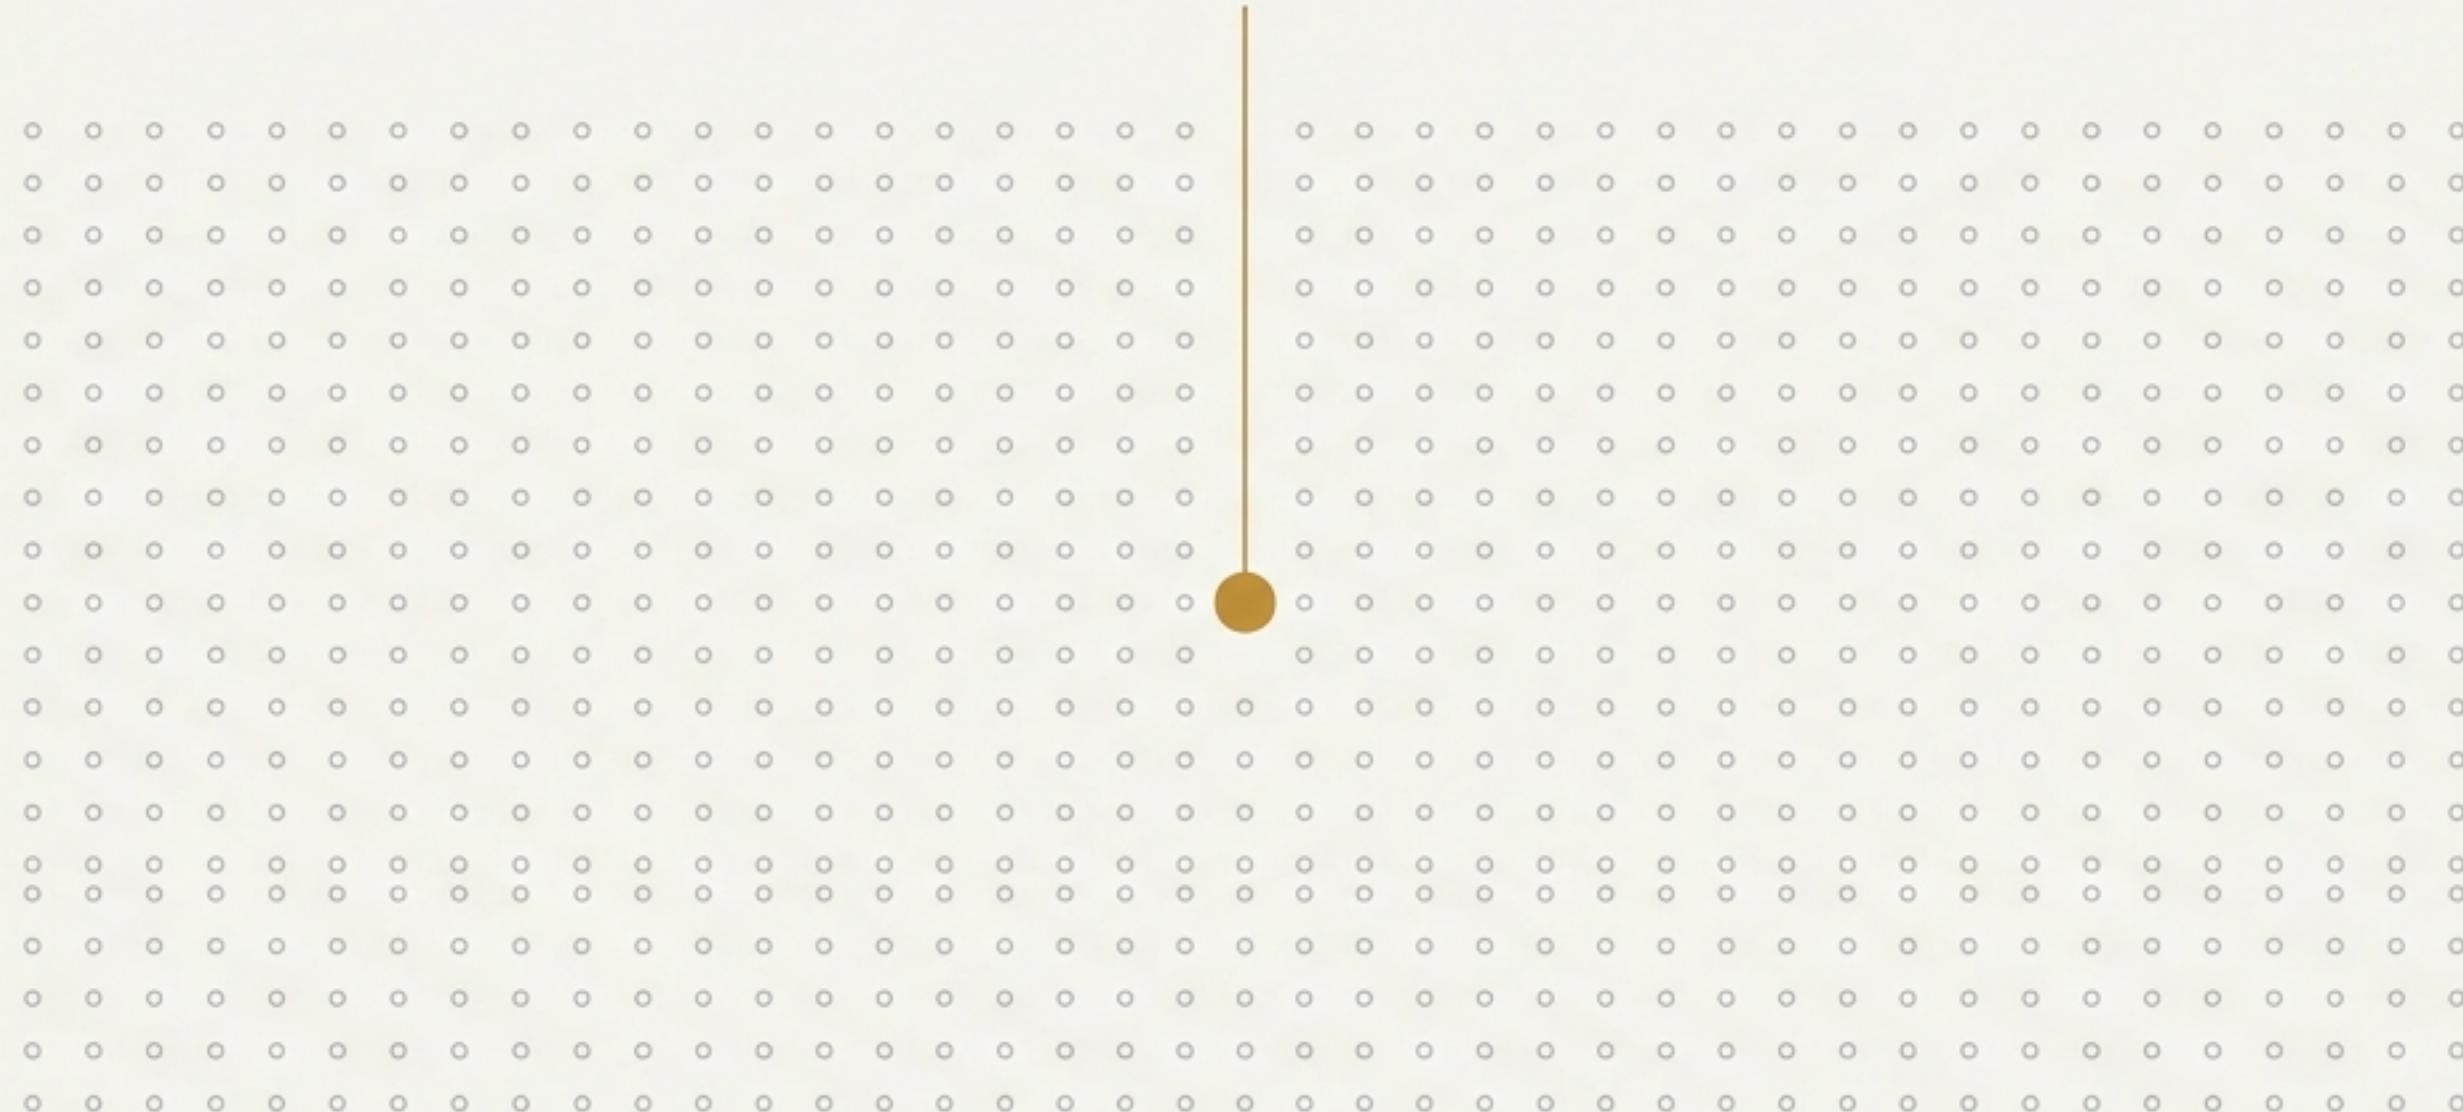


The Quest for Instant



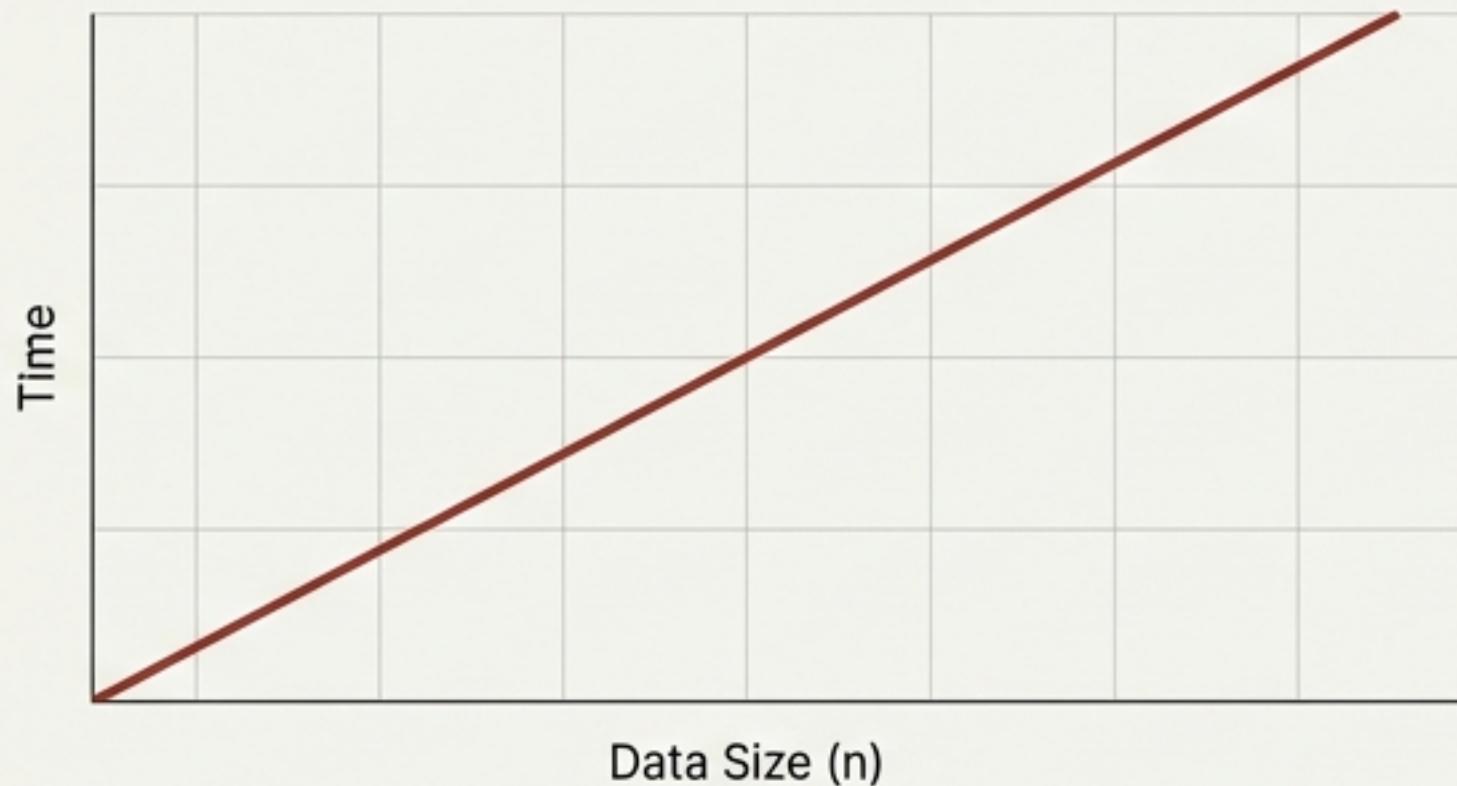
Mastering Data Retrieval with Hash Tables

The Fundamental Challenge: Finding Data Quickly

In a vast sea of data, finding a single piece of information can be slow. Imagine searching for one specific book in a library with no catalogue system.

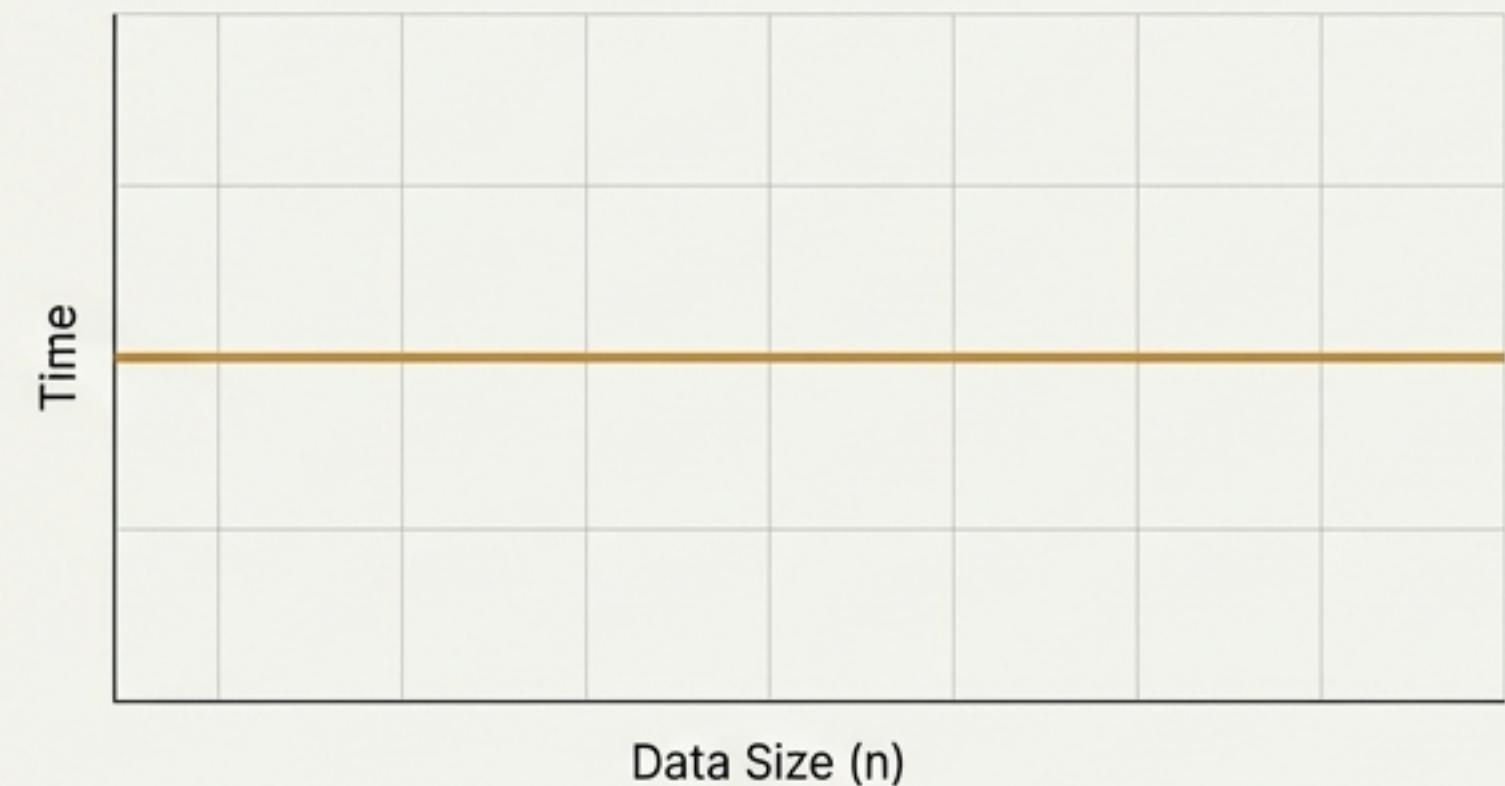
Linear Search: $O(n)$

To find an item, you must check every single entry one by one. The time taken grows directly with the amount of data.



The Goal: $O(1)$

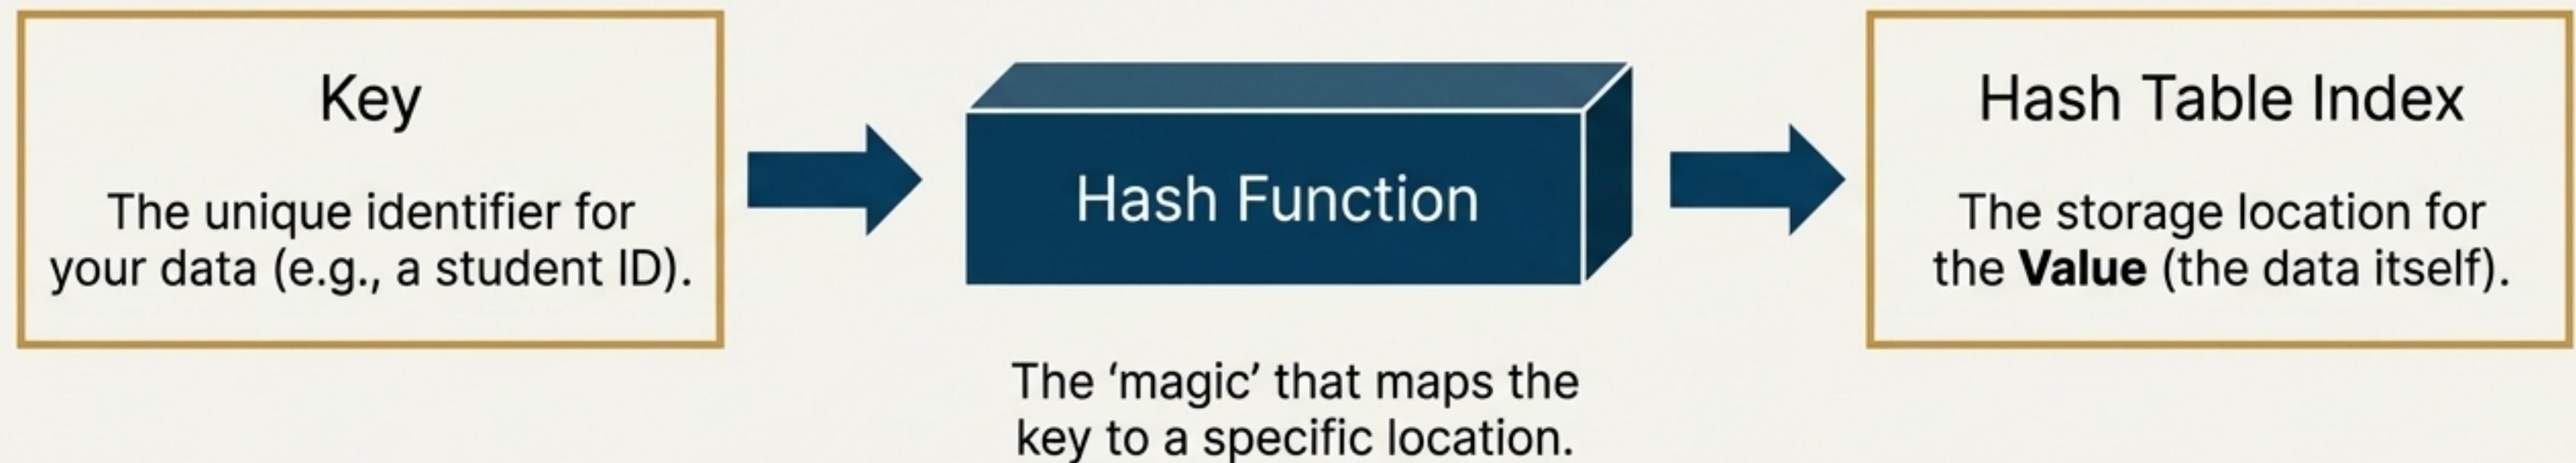
Accessing any item takes the same, constant amount of time, regardless of how much data you have. It's the equivalent of knowing the exact shelf and position of your book.



The Solution: The Hash Table

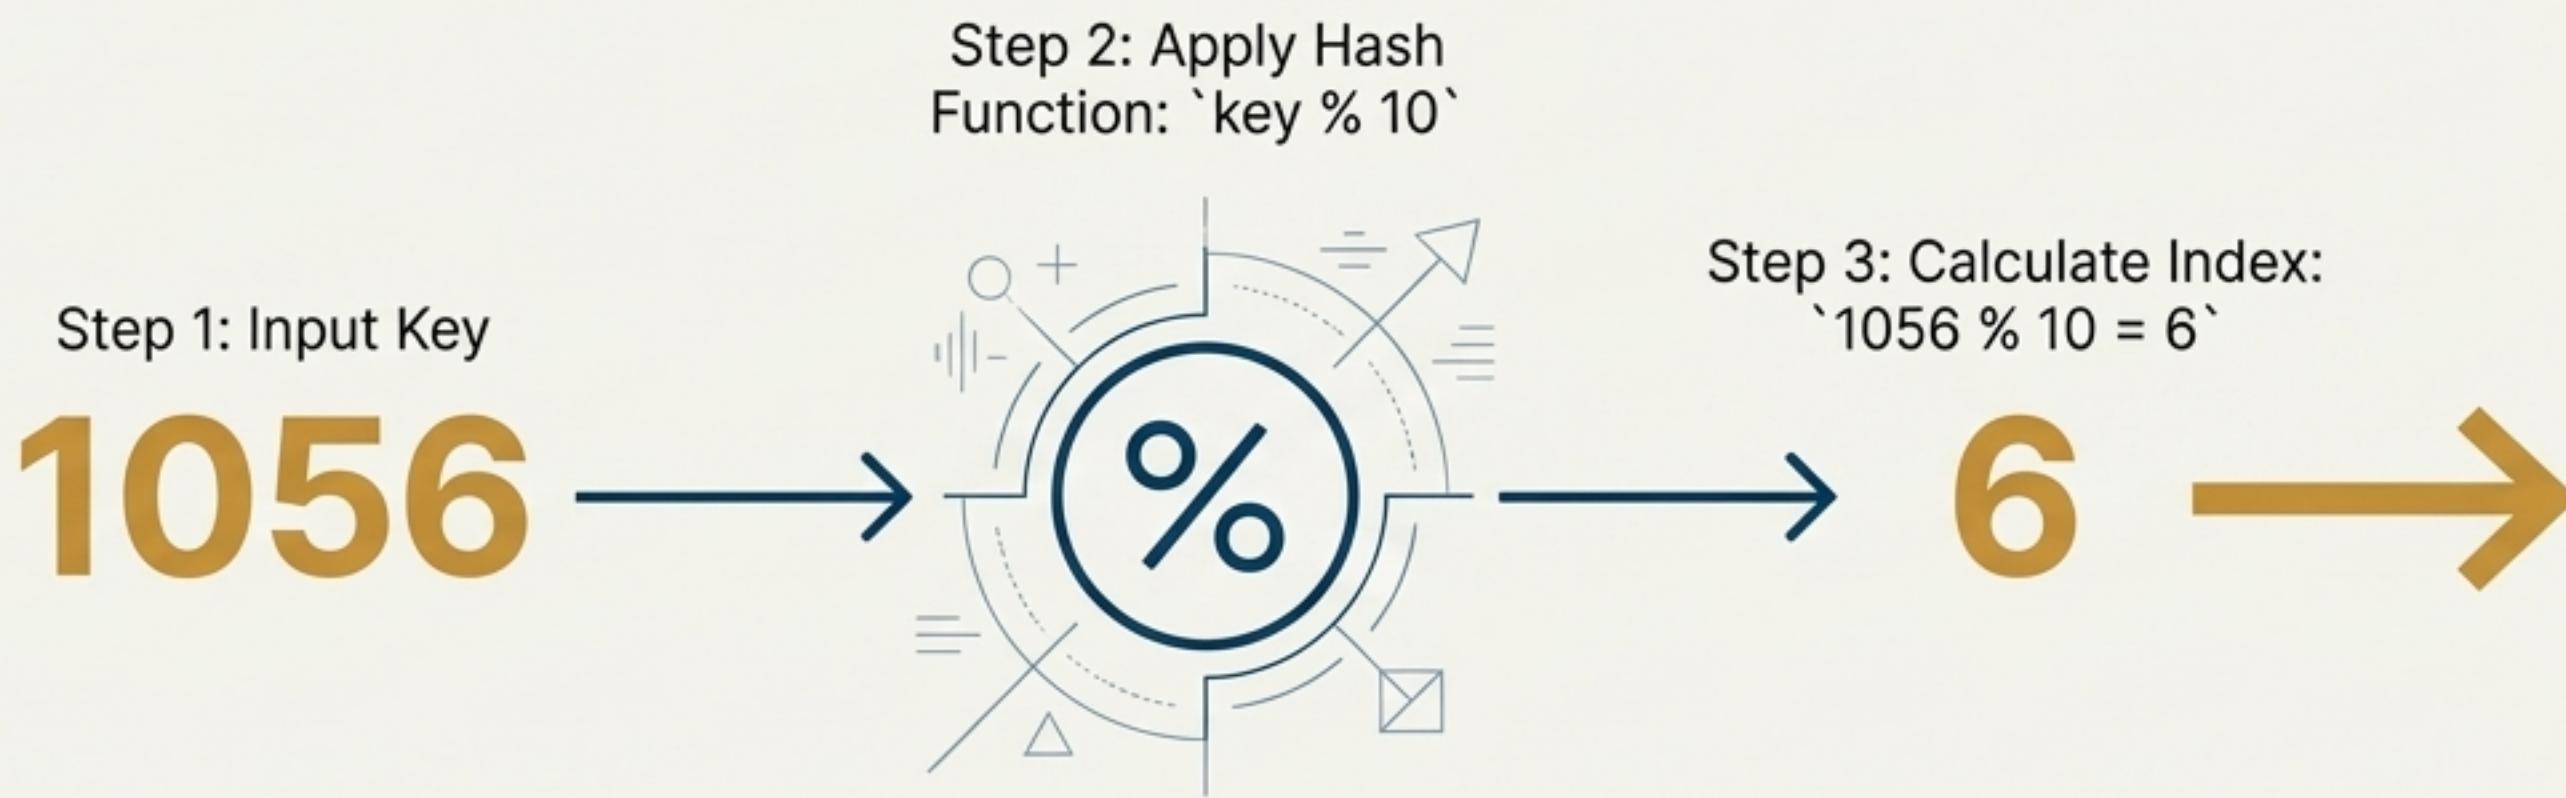
A Hash Table is a specialised data structure designed to achieve $O(1)$ average time complexity for insertion, deletion, and search operations.

It uses a clever mechanism to directly compute the location of your data, rather than searching for it.



From Key to Index: The Hash Function in Action

A hash function is a technique that converts a large key into a small, manageable index. A common and simple method is using the modulo operator.

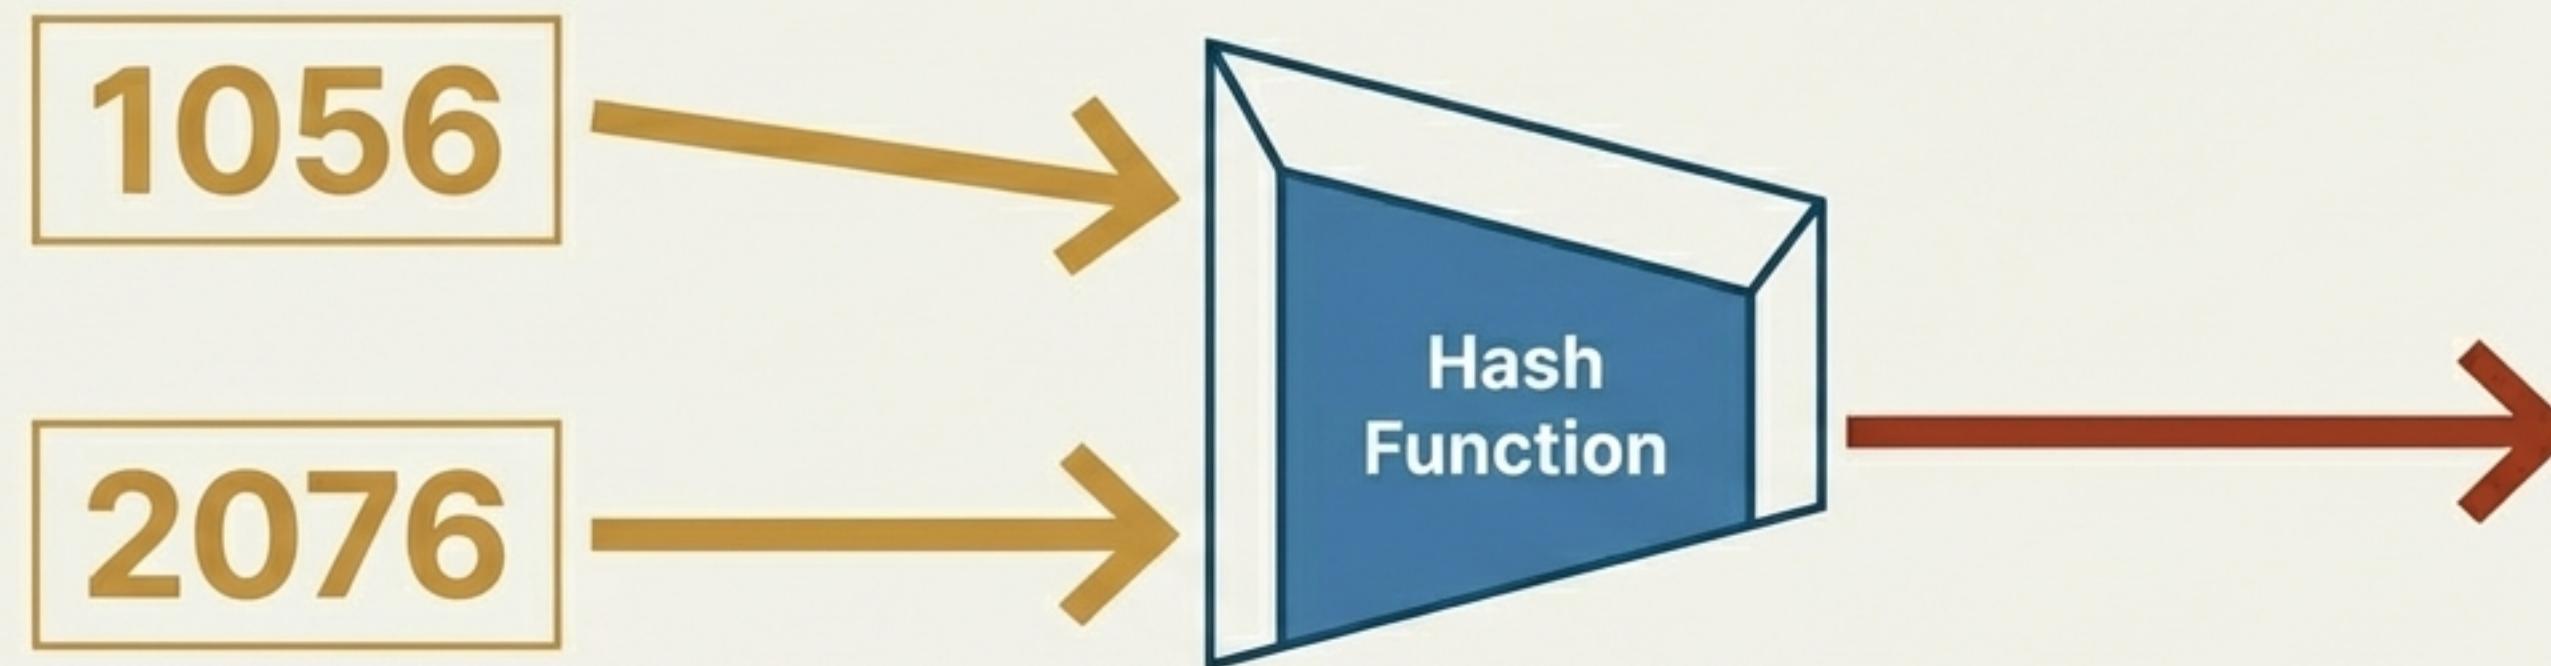


Step 4: Store Data

0	
1	
2	
3	
4	
5	
6	1056
7	
8	
9	

The Inevitable Conflict: When Keys Collide

A collision occurs when two different keys produce the exact same index after being processed by the hash function. The system's elegance is broken.



The Dilemma:** Both keys are trying to occupy the same single slot.
How do we store both?

0	
1	
2	
3	
4	
5	
6	1056 >< 2076
7	
8	
9	

Resolving the Conflict: Two Core Strategies

Engineers have developed two main techniques to handle collisions. Each has its own trade-offs in terms of memory, complexity, and performance.



Strategy 1: Chaining (External Hashing)

Philosophy: "If a slot is taken, build more room outside the table."

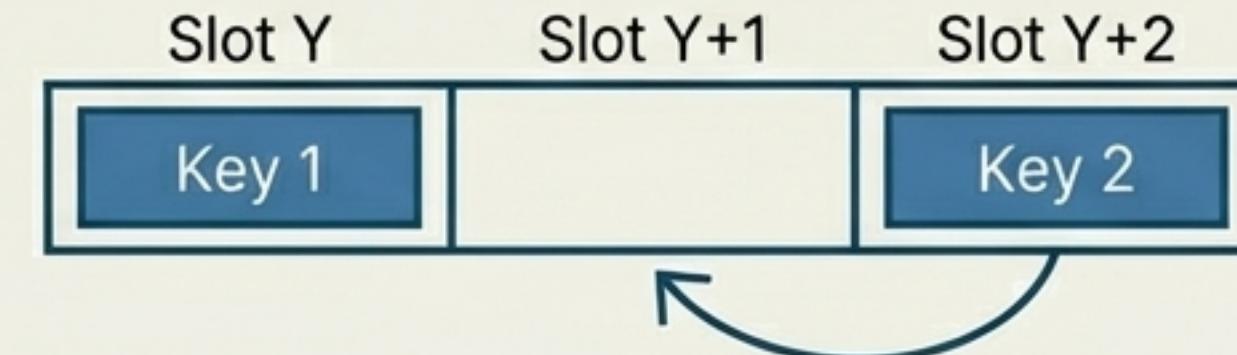
Mechanism: Each slot in the hash table acts as the head of a linked list. All keys that hash to the same index are stored in that index's list.



Strategy 2: Open Addressing (Closed Hashing)

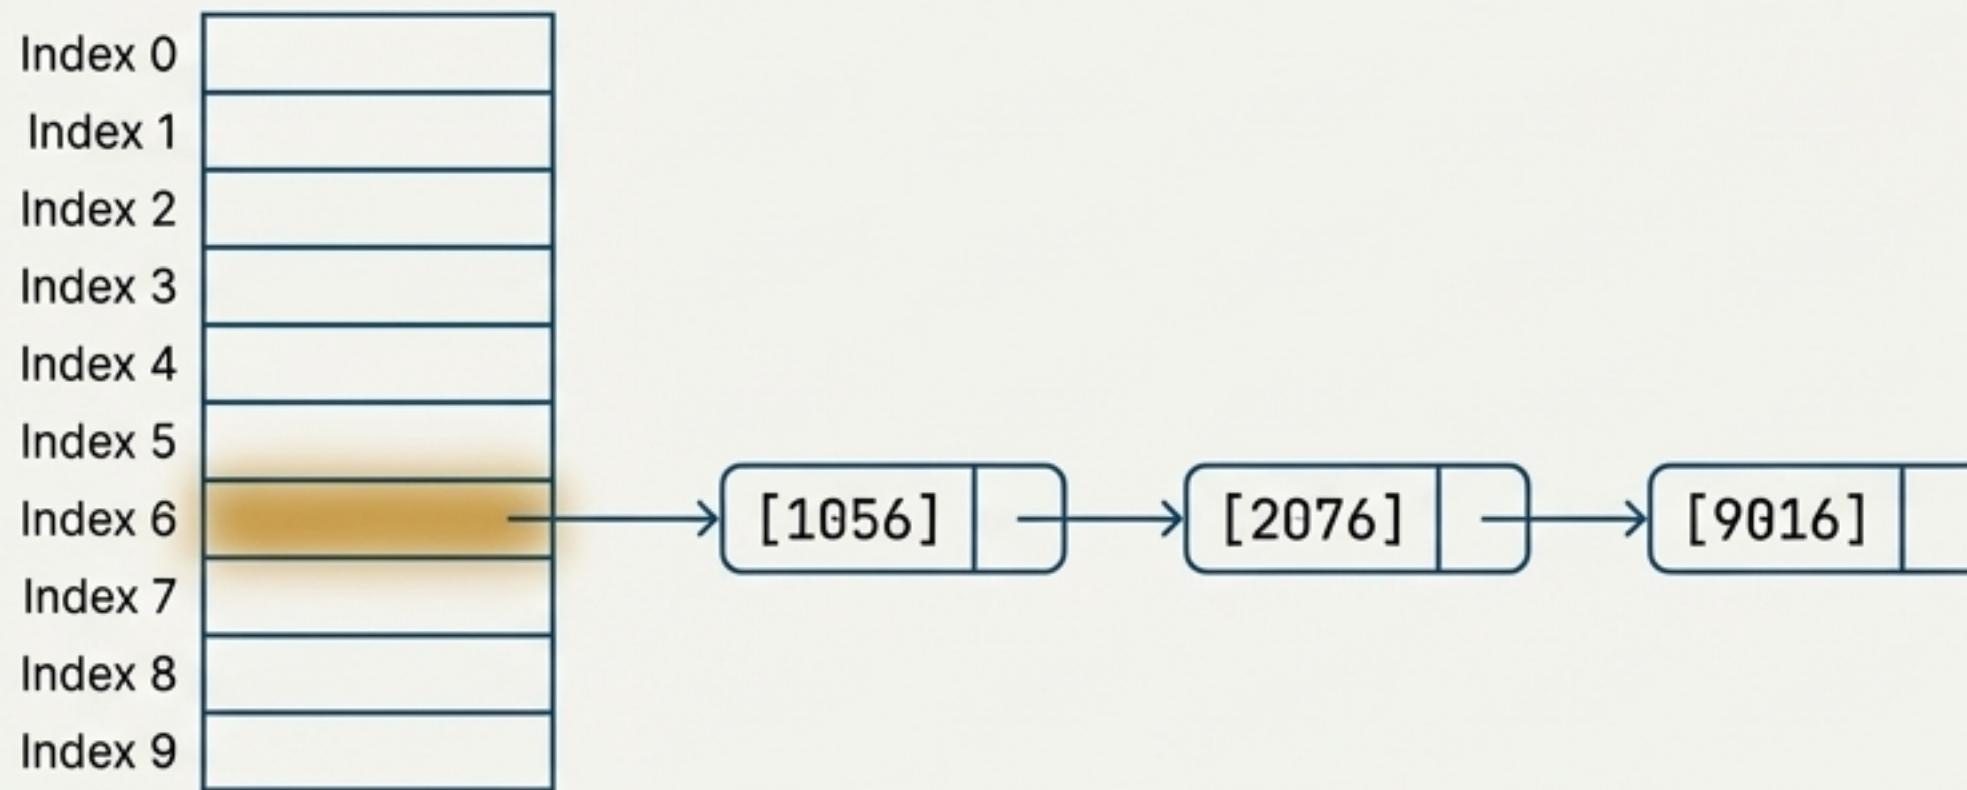
Philosophy: "If a slot is taken, cleverly find another empty slot inside the table."

Mechanism: All elements are stored within the table itself. When a collision occurs, we "probe" or search for the next available slot according to a set rule.



Path One: Chaining

In this method, each hash table index points to a data structure (typically a linked list) that can store multiple values. When a collision occurs, the new key is simply appended to the list at the collided index.



Advantages

- ✓ **Simple Implementation:** Relatively straightforward to code.
- ✓ **Handles Full Tables:** The table never truly becomes "full"; performance degrades as lists get longer, but it can always accept new data.
- ✓ **Good for Large Datasets:** Effective for storing an unknown or large number of elements.

Disadvantages

- ✗ **Extra Memory:** Requires additional memory for the pointers and list structures.
- ✗ **Performance Degradation:** If the load factor is high, linked lists can become very long, turning O(1) search into O(n) in the worst case.

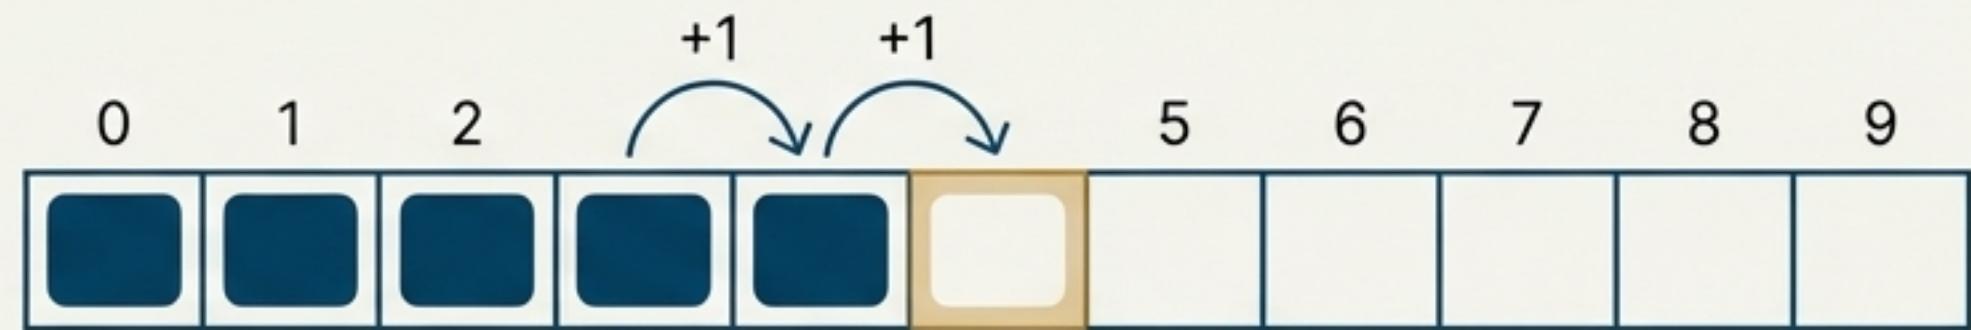
Path Two: Open Addressing

This method stores all elements directly inside the hash table array. When a collision happens at index $h(k)$, we probe for an alternative empty slot.

The Probing Sequence

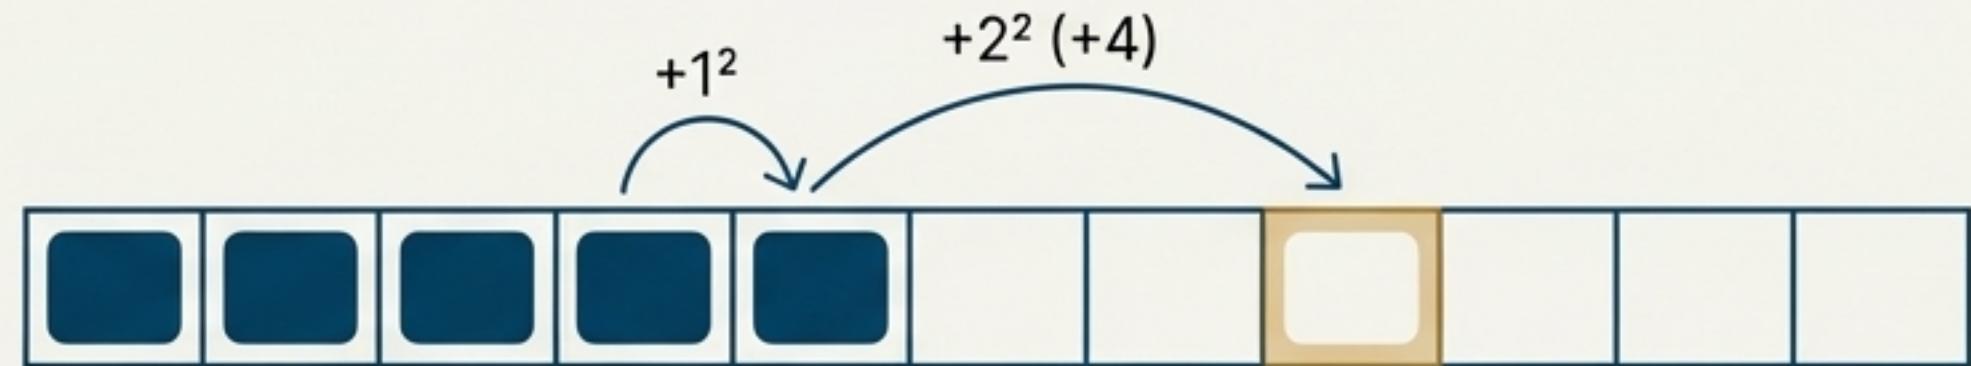
1. Linear Probing

If $h(k)$ is full, check $h(k)+1, h(k)+2, h(k)+3\dots$



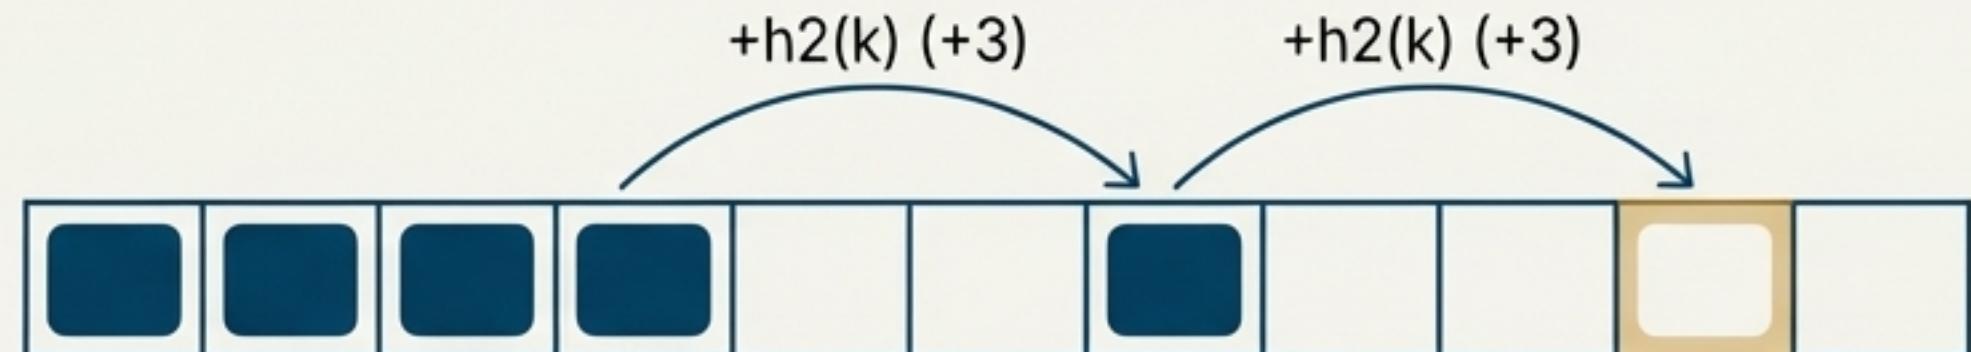
2. Quadratic Probing

To avoid clustering, check slots at increasing quadratic offsets:
 $h(k)+1^2, h(k)+2^2\dots$



3. Double Hashing

The most effective approach. Uses a second hash function to determine the step size: $h1(k) + i * h2(k)$.



The Blueprint Part I: Core Mechanics in C

```
#define SIZE 10
int hashTable[SIZE];

// Hash Function
int hashFunction(int key) {
    return key % SIZE;
}

// Insert using Linear Probing
void insert(int key) {
    int index = hashFunction(key);
    for (int i = 0; i < SIZE; i++) {
        int newIndex = (index + i) % SIZE;
        if (hashTable[newIndex] == -1 || hashTable[newIndex] == -2) {
            hashTable[newIndex] = key;
            return;
        }
    }
}
```

hashFunction(key): A clear and effective way to ensure the index is always within the array's bounds.

Linear Probing: This loop implements the probing mechanism. *i* represents the probe attempt number (0, 1, 2...).

Wrap-around Logic: The modulo operator ensures the search "wraps around" from the end of the table back to the beginning.

Finding a Slot: The crucial check. It inserts the key into the first available slot, which is either empty (`-1`) or previously deleted (`-2`).

The Blueprint Part 2: Searching and Deleting

```
// Search using Linear Probing
void search(int key) {
    int index = hashFunction(key);
    for (int i = 0; i < SIZE; i++) {
        int newIndex = (index + i) % SIZE;
        if (hashTable[newIndex] == key) return; // Found
        if (hashTable[newIndex] == -1) break; // Not found
    }
}

// Delete using Linear Probing
void deleteKey(int key) {
    // ... (probing loop as in search) ...
    if (hashTable[newIndex] == key) {
        hashTable[newIndex] = -2; // Mark as deleted
        return;
    }
}
```

The Deletion Problem

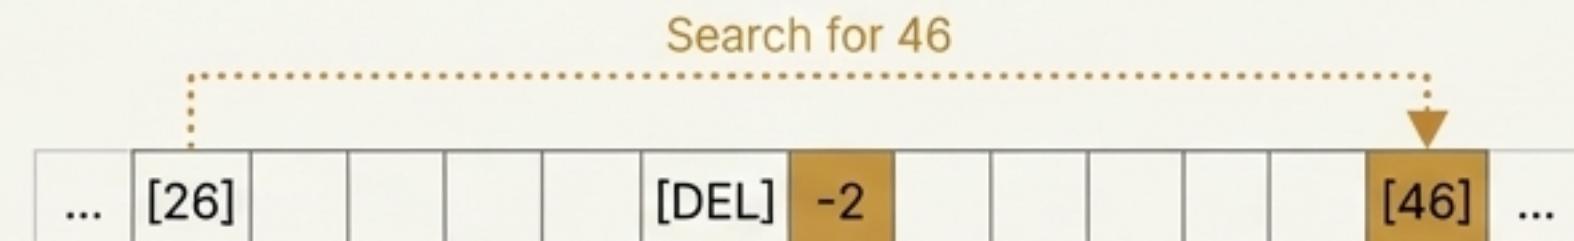
What happens if we just set a deleted slot back to -1 (empty)?

If we delete an item that was part of a collision chain, setting its slot to -1 would break the probe sequence. A subsequent search for an item further down the chain would hit the -1 and incorrectly conclude the item doesn't exist.



The Solution: The -2 Marker

We use a special value, -2, to mark the slot as "deleted but once occupied". The insert function can overwrite it, but the search function knows it must continue probing.



A World Transformed: The Impact of Hash Tables



Database Indexing

Maps unique IDs to entire data records for instant lookups.

Example: Student ID → Student Details.



Caching Systems

Maps a resource (like a URL) to its content for rapid retrieval.

Example: Browser cache maps URL → Webpage.



Compilers

Manages symbol tables, mapping variable and function names to their memory locations.

Typically resolved via Chaining.



Network Routers

Maps IP addresses to the next network hop, enabling high-speed packet routing.



File Systems

Maps a filename to its physical location on a storage device. *Used in systems like Linux EXT.*



Spell Checkers

Stores a dictionary of valid words for O(1) verification of a word's correctness.

The Quest for O(1): Key Takeaways

The Goal

We began with the need for constant-time, O(1), data access, the holy grail of efficiency.

The Hero

The Hash Table, which uses a Hash Function to directly map keys to storage indices.

The Conflict

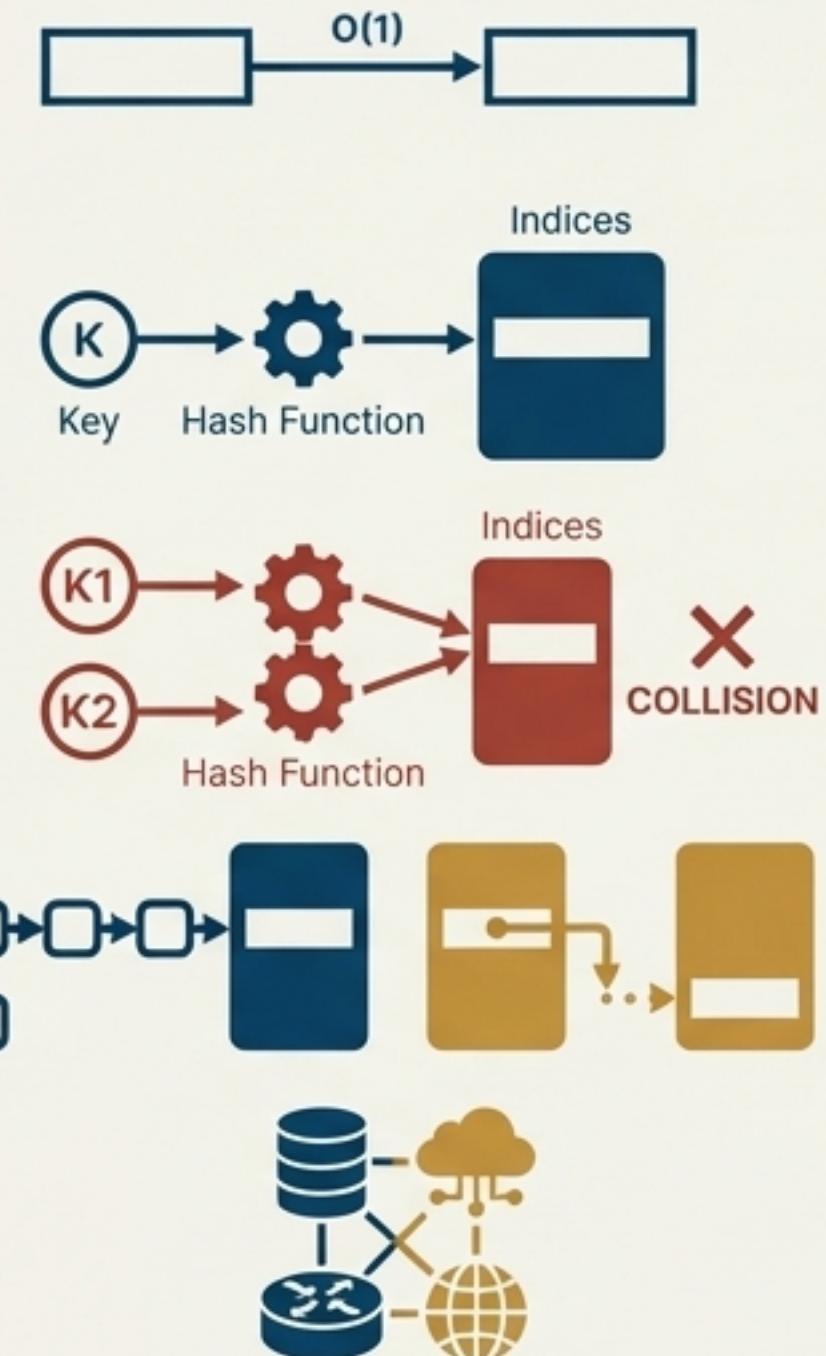
We encountered Collisions, where different keys map to the same index, challenging our system.

The Resolutions

We explored two powerful solutions: Chaining (building external lists) and Open Addressing (probing for internal space).

The Triumph

We saw how these principles power critical technologies, from databases to the internet itself.



Hash tables are not just a clever data structure; they are a fundamental building block of modern high-performance computing, making the 'instant' access we take for granted a daily reality.